



Validating DPDK performance on OpenShift

Updated Yesterday at 12:21 AM - English ▾

TABLE OF CONTENTS

[Validating DPDK performance on OpenShift](#)

[Traffic testing environment](#)

[Running the validation](#)

[Building the TRex container image](#)

[Deploying an OpenShift Container Platform cluster](#)

[Example performance profile](#)

[Example SR-IOV network policy](#)

[Example SR-IOV network](#)

[Defining a TRex pod](#)

[Create TRex configuration from template](#)

[Test configuration](#)

[Commands for TRex](#)

[Additional resources](#)

Validating DPDK performance on OpenShift


This article describes the build and the deployment of a traffic generating application inside a container. The traffic generator validates Data Plane Development Kit (DPDK) line rate performance on OpenShift Container Platform. It is used with the following network elements:

- Hardware networking

- Cloud elements
- Physical and virtual functions
- Test application

Traffic testing environment

The following diagram shows the components of a traffic-testing environment:



261_OpenShift_0722

- **Traffic generator:** An application that can generate high-volume packet traffic.
- **SR-IOV-supporting NIC:** A network interface card compatible with Single Root I/O Virtualization. The card runs a number of virtual functions on a physical interface.
- **Physical Function (PF):** A PCI Express (PCIe) function of a network adapter that supports the single root I/O virtualization (SR-IOV) interface.
- **Virtual Function (VF):** A lightweight PCIe function on a network adapter that supports Single Root I/O virtualization (SR-IOV). The VF is associated with the PCIe Physical Function (PF) on the network adapter, and represents a virtualized instance of the network adapter.
- **Switch:** Network switch. Nodes can also be connected back-to-back.
- **testpmd :** An example application included with DPDK. The `testpmd` application can be used to test the DPDK in a packet forwarding mode. `testpmd` is also an example of how to build a fully-fledged application using the DPDK SDK.
- **worker 0 and worker 1:** OpenShift Container Platform nodes.

Running the validation

For this stage of the development of the test environment, you must do the following:

1. Build the TRex container image
2. Deploy an OpenShift Container Platform cluster
3. Create TRex configuration from template
4. Define a TRex pod

Building the TRex container image

To build the TRex container image, run the following build script:

```
FROM quay.io/centos/centos:stream8

ARG TREX_VERSION=2.87
ENV TREX_VERSION ${TREX_VERSION}

# install requirements
RUN dnf install -y --nodocs git wget procps python3 vim python3-pip pciutils
gettext https://dl.fedoraproject.org/pub/epel/epel-release-latest-8.noarch.rpm &&
dnf clean all
RUN dnf install -y --nodocs hostname iproute net-tools ethtool nmap iputils perf
numactl sysstat htop rdma-core-devel libibverbs libibverbs-devel net-tools && dnf
clean all

# install trex server
WORKDIR /opt/
RUN wget --no-check-certificate https://trex-
tgn.cisco.com/trex/release/v${TREX_VERSION}.tar.gz && \
tar -xzf v${TREX_VERSION}.tar.gz && \
mv v${TREX_VERSION} trex && \
rm v${TREX_VERSION}.tar.gz

WORKDIR /opt/trex
```

Deploying an OpenShift Container Platform cluster

You must deploy and configure both the Node Tuning Operator and the SR-IOV Network Operator. See the OpenShift Container Platform 4.10 Documentation for more information. See also Additional Resources at the end of this article.

Example performance profile

The following code block illustrates a typical performance profile:

```
apiVersion: performance.openshift.io/v2
kind: PerformanceProfile
metadata:
  name: performance
spec:
  globallyDisableIrqLoadBalancing: true
  cpu:
    isolated: 21-51,73-103
    reserved: 0-20,52-72
  hugepages:
    defaultHugepagesSize: 1G
    pages:
      - count: 32
        size: 1G
  numa:
    topologyPolicy: "single-numa-node"
  nodeSelector:
    node-role.kubernetes.io/worker-cnf: ""
```

Description

- `isolated` : Defines the isolated CPUs for guaranteed workloads.
- `defaultHugepagesSize` : Defines the default `hugepages` size: typically set to 1G.
- `topologyPolicy` : Defines the Topology policy. The policy should always allocate from a single NUMA. If that is not possible, block the pod deployment.

Example SR-IOV network policy

The following code block illustrates a typical SR-IOV network policy:

```
apiVersion: sriovnetwork.openshift.io/v1
kind: SrioVNetworkNodePolicy
metadata:
  name: dpdk-nic-1
  namespace: openshift-sriov-network-operator
spec:
  deviceType: vfio-pci
  needVhostNet: true
  nicSelector:
    pfNames: ["ens3f0"]
  nodeSelector:
    node-role.kubernetes.io/worker-cnf: ""
  numVfs: 5
  priority: 99
  resourceName: dpdk_nic_1
---
apiVersion: sriovnetwork.openshift.io/v1
kind: SrioVNetworkNodePolicy
metadata:
  name: dpdk-nic-2
  namespace: openshift-sriov-network-operator
spec:
  deviceType: vfio-pci
  nicSelector:
    pfNames: ["ens3f1"]
  nodeSelector:
    node-role.kubernetes.io/worker-cnf: ""
  numVfs: 5
  priority: 99
  resourceName: dpdk_nic_2
```

Note: for Mellanix, use deviceType: netdevice and isRdma: True

Example SR-IOV network

The following code block illustrates a typical SR-IOV network:

```

---
apiVersion: sriovnetwork.openshift.io/v1
kind: SriovNetwork
metadata:
  name: dpdk-network-1-vlan
  namespace: openshift-sriov-network-operator
spec:
  ipam: '{"type": "host-local","ranges": [{"subnet": "10.0.1.0/24"}],"dataDir": "/run/my-orchestrator/container-ipam-state-1"}'
  networkNamespace: seba
  spoofChk: "on"
  trust: "on"
  vlan: 2004
  resourceName: dpdk_nic_1
---
apiVersion: sriovnetwork.openshift.io/v1
kind: SriovNetwork
metadata:
  name: dpdk-network-2
  namespace: openshift-sriov-network-operator
spec:
  ipam: '{"type": "host-local","ranges": [{"subnet": "10.0.2.0/24"}],"dataDir": "/run/my-orchestrator/container-ipam-state-2"}'
  networkNamespace: seba
  spoofChk: "on"
  trust: "on"
  resourceName: dpdk_nic_2

```

Note: Here, we are using `vlan` tag for the VFs as an example. This is not mandatory

Defining a TReX pod

The following `yaml` file defines a TReX pod:

```

---
apiVersion: v1
kind: Namespace
metadata:
  name: dpdk
---
apiVersion: v1
kind: ConfigMap
metadata:
  name: trex-info-for-config
  namespace: dpdk
data:
  PORT_BANDWIDTH_GB: "25"
---
apiVersion: v1
kind: ConfigMap
metadata:
  name: trex-config-template
data:
  trex_cfg.yaml : |
    - port_limit: 2
      version: 2
      interfaces:
        - "${PCIDEVICE_OPENSHIFT_IO_DPDK_NIC_1}"
        - "${PCIDEVICE_OPENSHIFT_IO_DPDK_NIC_2}"
      port_bandwidth_gb: ${PORT_BANDWIDTH_GB}
      port_info:
        - ip: 10.10.10.2
          default_gw: 10.10.10.1
        - ip: 10.10.20.2
          default_gw: 10.10.20.1
    platform:
      master_thread_id: $MASTER
      latency_thread_id: $LATENCY
      dual_if:
        - socket: ${SOCKET}
          threads: [${CPU}]
---
apiVersion: v1
kind: ConfigMap
metadata:
  name: trex-tests
  namespace: dpdk
data:
  testpmd.py : |
    from trex_stl_lib.api import *

    from testpmd_addr import *

    # Wild local MACs
    mac_localport0='50:00:00:00:00:01'

```

```

mac_localport1='50:00:00:00:00:02'

class STLS1(object):

    def __init__ (self):
        self.fsize = 64; # the size of the packet
        self.number = 0

    def create_stream (self, direction = 0):

        size = self.fsize - 4; # HW will add 4 bytes ethernet FCS
        dport = 1026 + self.number
        self.number = self.number + 1
        if direction == 0:
            base_pkt =
Ether(dst=mac_telco0,src=mac_localport0)/IP(src="16.0.0.1",dst=ip_telco0)/UDP(dport
=15,sport=1026)
        else:
            base_pkt =
Ether(dst=mac_telco1,src=mac_localport1)/IP(src="16.1.0.1",dst=ip_telco1)/UDP(dport
=16,sport=1026)
            #pad = max(0, size - len(base_pkt)) * 'x'
            pad = (60 - len(base_pkt)) * 'x'

        return STLStream(
            packet =
            STLPktBuilder(
                pkt = base_pkt / pad
            ),
            mode = STLTXCont()))

    def create_stats_stream (self, rate_pps = 1000, pgid = 7, direction = 0):

        size = self.fsize - 4; # HW will add 4 bytes ethernet FCS
        if direction == 0:
            base_pkt =
Ether(dst=mac_telco0,src=mac_localport0)/IP(src="17.0.0.1",dst=ip_telco0)/UDP(dport
=dport,sport=1026)
        else:
            base_pkt =
Ether(dst=mac_telco1,src=mac_localport1)/IP(src="17.1.0.1",dst=ip_telco1)/UDP(dport
=dport,sport=1026)
            pad = max(0, size - len(base_pkt)) * 'x'

        return STLStream(
            packet =
            STLPktBuilder(
                pkt = base_pkt / pad
            ),
            mode = STLTXCont(pps = rate_pps),
            flow_stats = STLFlowLatencyStats(pg_id = pgid))

```

```

#flow_stats = STLFlowStats(pg_id = pgid))

def get_streams (self, direction = 0, **kwargs):
    # create multiple streams, one stream per core...
    s = []
    for i in range(14):
        s.append(self.create_stream(direction = direction))
    #if direction == 0:
    #    s.append(self.create_stats_stream(rate_pps=1000, pgid=10,
direction = direction))
    #else:
    #    s.append(self.create_stats_stream(rate_pps=1000, pgid=11,
direction = direction))

    return s

# dynamic load - used for trex console or simulator
def register():
    return STLS1()

testpmd_addr.py: |
# wild second XL710 mac
mac_telco0 = '60:00:00:00:00:01'
# we don't care of the IP in this phase
ip_telco0  = '10.0.0.1'
# wild first XL710 mac
mac_telco1 = '60:00:00:00:00:02'
ip_telco1  = '10.1.1.1'

---
apiVersion: v1
kind: Pod
metadata:
  annotations:
    k8s.v1.cni.cncf.io/networks: '[{"name": "dpdk-network-1-vlan", "mac": "50:00:00:00:00:01", "namespace": "dpdk"}, {"name": "dpdk-network-2", "mac": "50:00:00:00:00:02", "namespace": "dpdk"}]'
  cpu-load-balancing.crio.io: "disable"
  cpu-quota.crio.io: "disable"
  irq-load-balancing.crio.io: "disable"
  labels:
    app: trex
  name: trex

```

```
namespace: dpdk
spec:
  runtimeClassName: performance-performance
  affinity:
    podAntiAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        - labelSelector:
            matchExpressions:
              - key: app
                operator: In
                values:
                  - dpdk
      topologyKey: kubernetes.io/hostname
  containers:
    - command:
        - /bin/bash
        - -c
        - sleep INF
    image: <TREX-IMAGE>
    imagePullPolicy: Always
    name: trex
    envFrom:
      - configMapRef:
          name: trex-info-for-config
  resources:
    limits:
      cpu: "16"
      hugepages-1Gi: 8Gi
      memory: 1Gi
    requests:
      cpu: "16"
      hugepages-1Gi: 8Gi
      memory: 1Gi
  securityContext:
    capabilities:
      add:
        - IPC_LOCK
        - SYS_RESOURCE
        - NET_RAW
        - NET_ADMIN
    runAsUser: 0
  volumeMounts:
    - name: trex-config-template
      mountPath: /opt/templates/
    - name: trex-tests
      mountPath: /opt/tests/
    - mountPath: /mnt/huge
      name: hugepages
    - name: modules
      mountPath: /lib/modules
```

```

terminationGracePeriodSeconds: 5
volumes:
  - name: modules
    hostPath:
      path: /lib/modules
  - configMap:
      name: trex-info-for-config
    name: trex-info-for-config
  - name: trex-config-template
    configMap:
      name: trex-config-template
  - name: trex-tests
    configMap:
      name: trex-tests
  - emptyDir:
      medium: HugePages
    name: hugepages

```

Description

- PCIDEVICE_OPENSHIFT_IO_DPDK_NIC_1 : The PCI deviceID for the first network.
- PCIDEVICE_OPENSHIFT_IO_DPDK_NIC_2 : The PCI deviceID for the second network.
- Trex-config-template : The template configuration. Create this before starting the `trex` process.
- TREX-IMAGE : the TRex image built before the deployment

Create TRex configuration from template

Access the TRex pod with `oc -n dpdk rsh trex` the copy the template file `cp /opt/templates/trex_cfg.yaml /etc/trex_cfg.yaml`
and edit the file using the following parameters

```

port_limit: 2
version: 2
interfaces: ["{PCIDEVICE_1}","{PCIDEVICE_2}"]
port_bandwidth_gb: {PORT_BANDWIDTH_GB}
port_info:
  - ip: 10.10.10.2
    default_gw: 10.10.10.1
  - ip: 10.10.20.2
    default_gw: 10.10.20.1
platform:
  master_thread_id: {MASTER}
  latency_thread_id: {LATENCY}
  dual_if:
    - socket: {SOCKET}
      threads: [{CPUS}]

```

replace the parameters with the following configuration

PCIDEVICE_1: check the PCI address that was allocated for the trex pod running `env | grep PCIDEVICE_OPENSHIFT_IO`

PCIDEVICE_2 : check the PCI address that was allocated for the trex pod running `env | grep PCIDEVICE_OPENSHIFT_IO`

PORt_BANDWIDTH_GB: the port speed in GB (example: 25)

MASTER: select the master/UI CPU. check the available CPUs for the pod using `cat /sys/fs/cgroup/cpuset/cpuset.cpus`. Check also for sibling CPUs example for cpu0 (`cat /sys/devices/system/cpu/cpu0/topology/core_cpus_list`)

LATENCY: select the latency measurement CPU (better to use the sibling from the MASTER CPU selected)

SOCKET: select the Numa where all the allocated CPUs are (you can check using the `lspcu | grep`

CPU: add all the CPUs from `cat /sys/fs/cgroup/cpuset/cpuset.cpus` excluding the ones used by MASTER and LATENCY

Output Example:

```
- port_limit: 2
  version: 2
  interfaces:
    - "0000:d8:02.1"
    - "0000:d8:0a.3"
  port_bandwidth_gb: 25
  port_info:
    - ip: 10.10.10.2
      default_gw: 10.10.10.1
    - ip: 10.10.20.2
      default_gw: 10.10.20.1
  platform:
    master_thread_id: 21
    latency_thread_id: 23
  dual_if:
    - socket: 1
      threads: [25,27,29,31,33,35,73,75,77,79,81,83,85,87]
```

Then you can launch TRex with the following command:

```
./t-rex-64 --no-ofed-check --no-hw-flow-stat -i -c 14
```

Test configuration

The test is injected into the container using a `configmap`. Using this method, you do not need to recreate the pod if you change the code. Apply the `configmap` again and `k8s` will mount the new files.

The following is an example `testpmd.py`:

```

from trex_stl_lib.api import *

from testpmd_addr import *

# Wild local MACs
mac_localport0='50:00:00:00:00:01'
mac_localport1='50:00:00:00:00:02'

class STLS1(object):

    def __init__ (self):
        self.fsize = 64; # the size of the packet
        self.number = 0

    def create_stream (self, direction = 0):

        size = self.fsize - 4; # HW will add 4 bytes ethernet FCS
        dport = 1026 + self.number
        self.number = self.number + 1
        if direction == 0:
            base_pkt =
Ether(dst=mac_telco0,src=mac_localport0)/IP(src="16.0.0.1",dst=ip_telco0)/UDP(dport=15,sport=1026)
        else:
            base_pkt =
Ether(dst=mac_telco1,src=mac_localport1)/IP(src="16.1.0.1",dst=ip_telco1)/UDP(dport=16,sport=1026)
        #pad = max(0, size - len(base_pkt)) * 'x'
        pad = (60 - len(base_pkt)) * 'x'

        return STLStream(
            packet =
            STLPktBuilder(
                pkt = base_pkt / pad
            ),
            mode = STLTXCont()))

    def create_stats_stream (self, rate_pps = 1000, pgid = 7, direction = 0):

        size = self.fsize - 4; # HW will add 4 bytes ethernet FCS
        if direction == 0:
            base_pkt =
Ether(dst=mac_telco0,src=mac_localport0)/IP(src="17.0.0.1",dst=ip_telco0)/UDP(dport=dport,sport=1026)
        else:
            base_pkt =
Ether(dst=mac_telco1,src=mac_localport1)/IP(src="17.1.0.1",dst=ip_telco1)/UDP(dport=dport,sport=1026)
        pad = max(0, size - len(base_pkt)) * 'x'

        return STLStream(

```

```

packet =
STLPktBuilder(
pkt = base_pkt / pad
),
mode = STLTXCont(pps = rate_pps),
flow_stats = STLFlowLatencyStats(pg_id = pgid))

def get_streams (self, direction = 0, **kwargs):
# create multiple streams, one stream per core...
s = []
for i in range(14):
s.append(self.create_stream(direction = direction))
return s

# dynamic load - used for trex console or simulator
def register():
return STLS1()

```

The following is an example `testpmd_addr.py`:

```

# wild second XL710 mac
mac_telco0 = '60:00:00:00:00:01'
# we don't care of the IP in this phase
ip_telco0  = '10.0.0.1'
# wild first XL710 mac
mac_telco1 = '60:00:00:00:00:02'
ip_telco1  = '10.1.1.1'

```

Commands for TRex

From a new terminal access the TRex pod `oc -n dpdk rsh trex`, and start the connection to TRex. You can then load the traffic generator stream.

```

./trex-console
> tui
> # start -f /opt/tests/testpmd.py -m <number-of-packets> -p <ports to use>
> start -f /opt/tests/testpmd.py -m 24mpps -p 0
> stop -a

```

This code example runs the traffic generator, sending 24 million packets per second using port 0 and receiving from port 1:

Example output

Global Statistics

```
connection : localhost, Port 4501  total_tx_L2 : 12.76 Gbps
version   : STL @ v2.87total_tx_L1 : 16.74 Gbps
cpu_util.: 29.65% @ 14 cores (14 per dual port)  total_rx : 12.09 Gbps
rx_cpu_util. : 0.0% / 0 pps  total_pps: 24.91 Mpps
async_util. : 0.02% / 10.48 Kbps drop_rate: 0 bps
total_cps. : 0 cps  queue_full : 0 pkts
```

Port Statistics

port	0	1	total
owner	root	root	
link	UP	UP	
state	IDLE	TRANSMITTING	
speed	25 Gb/s	25 Gb/s	
CPU util.	0.0%	29.65%	
--			
Tx bps L2	0 bps	12.76 Gbps	12.76 Gbps
Tx bps L1	0 bps	16.74 Gbps	16.74 Gbps
Tx pps	0 pps	24.91 Mpps	24.91 Mpps
Line Util.	0 %	66.97 %	
--			
Rx bps	12.09 Gbps	0 bps	12.09 Gbps
Rx pps	20.99 Mpps	0 pps	20.99 Mpps
--			
opackets	2654147973	1223919476	3878067449
ipackets	1030488553	2012583907	3043072460
obytes	180482062450	78330844946	258812907396
ibytes	74195178964	136855711844	211050890808
tx-pkts	2.65 Gpkts	1.22 Gpkts	3.88 Gpkts
rx-pkts	1.03 Gpkts	2.01 Gpkts	3.04 Gpkts
tx-bytes	180.48 GB	78.33 GB	258.81 GB
rx-bytes	74.2 GB	136.86 GB	211.05 GB
--			
oerrors	0	0	0
ierrors	0	0	0

On the other side of the traffic testing environment is a machine that can process the packets and send them back to the TRex machine on port 1. Possible usages are `testpmd`, `vpp` and `ovs-dpdk`.

Additional resources

- About the Performance Profile Creator
- Adjusting the NIC queues with the performance profile

- Provisioning a worker with real-time capabilities
- Installing the SR-IOV Network Operator
- SR-IOV network node configuration object
- Dynamic IP address assignment configuration with Whereabouts

SBR Networking

Product(s) Red Hat OpenShift Container Platform

Category Developer

Article Type General

Was this helpful?

YES

NO

People who viewed this article also viewed

**How to
troubleshoot OVS
DPDK performance
issue**

Solution - May 18,
2024


**How to
troubleshoot OVS
DPDK performance
issue**

Solution - Jun 14,
2024

**OpenShift on
OpenStack with
DPDK**

Solution - Jun 14,
2024

Comments



RED HAT

NEWBIE

Add comment

[Formatting Help](#)

Send notifications to content followers

Mark comment as private

Submit



RED HAT

ACTIVE CONTRIB...

148 Points

Private Comment May 22, 2023 1:27 PM

[Anil Dhingra](#)

was not able to start Trex - so started with --no-scapy-server

```
sh-4.4# ./t-rex-64 --no-ofed-check --no-hw-flow-stat -i -c 3 --no-scapy-server
```

Try troubleshoot why scapy not starting - assuming without this option may not be able to run .py profiles

[Reply Privately](#)



RED HAT

ACTIVE CONTRIB...

148 Points

Private Comment May 25, 2023 7:30 AM

[Anil Dhingra](#)

Worth to highlight why below parameter is required , users are facing packet drop and multiple performance related cases from partners/vendors

cpu-quota.crio.io: "disable" runtimeClassName: performance-performance

<https://issues.redhat.com/browse/RFE-3758>

 [Reply Privately](#)



Quick Links

Help

Site Info

Related Sites

About

[Red Hat Subscription Value](#)

[About Red Hat](#)

[Red Hat Jobs](#)



[About Red Hat](#)

[Jobs](#)

[Events](#)

[Locations](#)

[Contact Red Hat](#)

[Red Hat Blog](#)

[Inclusion at Red Hat](#)

[Cool Stuff Store](#)

[Red Hat Summit](#)

Copyright © 2025 Red Hat, Inc.

[Privacy statement](#)

[Terms of use](#)

[All policies and guidelines](#)

[Digital accessibility](#)

[Cookie preferences](#)